

Getting Started with R

Sven Otto

February 26, 2024

Table of contents

Welcome	3
Why R?	3
Matrix algebra	4
Accompanying R scripts	4
Comments	4
 1 Base R	 5
1.1 Short Glossary	5
1.2 First Steps	5
1.3 Vectors and functions	6
1.4 Further Data Objects	9
1.4.1 The matrix	9
1.4.2 The list	10
1.4.3 The data frame	11
1.4.4 The ts object	13
 2 Packages	 16
2.1 The xts package	16
2.2 Data packages	17
2.3 The tidyverse	20

Welcome

This tutorial aims to serve as an introduction to the software package R. Other excellent and much more exhaustive tutorials can be found at the following links:

- An interactive R-package for learning R: [swirl](#) (highly recommended for beginners).
- Interactive R courses at [Datacamp](#) and [Coursera](#) (free, but registration required).
- Video series by Nick Huntington-Klein: [Introduction to R for Economists](#).
- [The official introduction](#) and reference cards for [basic R](#) and [time series analysis](#).
- Some excellent books:
 - [Hands-On Programming with R](#) (for absolute beginners)
 - [R for Data Science](#) (R and the tidyverse)
 - [Advanced R](#) (improve your programming skills)
 - [R Cookbook](#) (proven recipes for data analysis)
 - [Forecasting: Principles and Practice](#) (time series analysis in R)
 - [R Packages](#) (write your own R package)
 - [HappyGitWithR](#) (version control with RStudio)

Why R?

- R is **free** of charge. On the R project webpage cran.r-project.org, you can download R for Windows, Mac OS, or Linux. Windows users can also directly follow this link: cran.r-project.org/bin/windows/base/
- You can use R via a terminal or install an IDE, which is much more convenient. The celebrated IDE **RStudio** for R is also **free** of charge. Download RStudio here: posit.co/download/rstudio-desktop/. Make sure that you install R before installing RStudio.
- Within RStudio, you can use **Quarto**, which provides an authoring framework to export your R code/outputs/plots together with LaTeX formulas and text as a PDF file or website in an appealing way. Have a look [here](#). This website is also built with Quarto. You may want to use Quarto for your assignments, term papers, or thesis.
- R is equipped with one of the most flexible and powerful graphics routines available anywhere. Check out these repositories with examples of appealing and informative R graphs: [Clean Graphs](#), [R Graph Catalog](#), [Publication Ready Plots](#).

- One of the best features of R are the large number of contributed packages from the statistical community. You find R packages for almost any statistical method out there and many statisticians provide R packages to accompany their research.
- R is the de-facto standard for statistical science.

Matrix algebra

R is a matrix-based programming language. Matrix algebra provides an efficient framework for analyzing and implementing econometric methods. To refresh your matrix algebra skills and to learn how to use it in R, please check out my [Crash Course on Matrix Algebra in R](#).

Accompanying R scripts

All R codes of the different sections can be found here:

- [rintro-sec1.R](#).
- [rintro-sec2.R](#).

Comments

Feedback is welcome. If you notice any typos or issues, please report them on [GitHub](#) or email me at sven.otto@uni-koeln.de.

1 Base R

1.1 Short Glossary

Let's start the tutorial with a (very) short glossary:

- **Console:** The thing with the `>` sign at the beginning.
- **Script file:** An ordinary text file with suffix `.R`. For instance, `yourfilename.R`.
- **Working directory:** The file directory you are working in. If no directory is explicitly specified when loading data, then R assumes that the data is located in the *working directory*. Useful commands: with `getwd()`, you get the location of your current working directory, and `setwd()` allows you to set a new location for it.
- **Workspace:** This is a hidden file (stored in the working directory as `.RData`) where all objects you use (e.g., data, matrices, vectors, variables, functions, etc.) are stored. When you close RStudio, you will be asked if you want to save or delete the session's *workspace*. If you save it, it will be loaded automatically with the next R session, provided you start R in the corresponding working directory. Useful commands: `ls()` shows all elements in our current workspace, and `rm(list=ls())` deletes all elements in our current workspace.

1.2 First Steps

A good idea is to use a script file like **myscript.R** to store your R commands. You can send single lines or marked areas of your R code to the console by pressing the **CTRL+RETURN** (STRG+ENTER) keys.

To start with baby steps, we do some simple calculations:

```
2+2 # addition
```

```
[1] 4
```

```
2*2 # multiplication
```

```
[1] 4
```

```
2/2
```

```
[1] 1
```

```
2-2
```

```
[1] 0
```

```
2^3 # exponentiate
```

```
[1] 8
```

Note: Anything written after the `#` sign will be ignored by R, which is very useful for commenting on your code.

The **assignment operator** `<-` will be your most often-used tool. Here is an example of creating a **scalar** variable:

```
x <- 4  
x
```

```
[1] 4
```

```
4 -> x # possible but unusual  
x
```

```
[1] 4
```

```
x = 4  
x
```

```
[1] 4
```

Note: The R community loves the `<-` assignment operator. Alternatively, you can use the `=` operator.

1.3 Vectors and functions

And now a more interesting object - a **vector**:

```
y = c(2,7,4,1)
y
```

```
[1] 2 7 4 1
```

The command `ls()` shows the total content of your current workspace, and the command `rm(list=ls())` deletes all elements of your current workspace:

```
ls()
```

```
[1] "has_annotations" "x" "y"
```

```
rm(list=ls())
ls()
```

```
character(0)
```

Note: RStudio's **Environment** pane also lists all the elements in your current workspace. That is, the command `ls()` becomes a bit obsolete when working with RStudio.

Let's try how we can compute with vectors and scalars in R.

```
x = 4
y = c(2,7,4,1)

x*y # each element in the vector y is multiplied by the scalar x
```

```
[1] 8 28 16 4
```

```
y*y # a term-by-term product of the elements in y
```

```
[1] 4 49 16 1
```

The term-by-term execution, as in the above example, `y*y`, is a main strength of R. We can conduct many operations **vector-wisely**:

```
y^2
```

```
[1]  4 49 16  1
```

```
log(y)
```

```
[1] 0.6931472 1.9459101 1.3862944 0.0000000
```

```
exp(y)
```

```
[1]  7.389056 1096.633158  54.598150  2.718282
```

```
y-mean(y)
```

```
[1] -1.5  3.5  0.5 -2.5
```

```
(y-mean(y))/sd(y) # standardization
```

```
[1] -0.5669467  1.3228757  0.1889822 -0.9449112
```

Element-wise operations are a central characteristic of matrix-based languages like R (or Matlab). Other programming languages often have to use **loops** instead:

```
N = length(y)
1:N

y.sq = rep(0,N)
y.sq

for(i in 1:N){
  y.sq[i] = y[i]^2
  if(i == N){
    print(y.sq)
  }
}
```


The `for()`-loop is the most common loop, but there is also a `while()`-loop and a `repeat()`-loop. However, loops in R can be relatively slow. Therefore, try to avoid them!

Useful commands to produce **sequences** of numbers:

```
1:10
-10:10
?seq # Help for the seq()-function
seq(from=1, to=100, by=7) # sequence generation
rep(0,10) # replicate elements
```

The `[]`-operator **selects elements** of vectors:

```
y[c(2,4)]
```

```
[1] 7 1
```

Element selections can be made on a more **logical** basis, too. For example, if you want only the elements of the vector `y` that are strictly greater than 2:

```
y[y>2]
```

```
[1] 7 4
```

```
# Note that this gives you a boolean vector:
y>2
```

```
[1] FALSE TRUE TRUE FALSE
```

Note: Logical operations return so-called **boolean** objects, i.e., a `TRUE` or a `FALSE`. For instance, if we ask R whether `1>2`, we get the answer `FALSE`.

1.4 Further Data Objects

Besides the classical data objects like scalars and vectors, there are many other objects in R:

1.4.1 The matrix

A **matrix** is a rectangular array of numbers.

```
mymatrix = matrix(data=1:16, nrow=4, ncol=4)
mymatrix
```

```
      [,1] [,2] [,3] [,4]
[1,]     1     5     9    13
[2,]     2     6    10    14
[3,]     3     7    11    15
[4,]     4     8    12    16
```

Matrices are extremely useful for theoretically analyzing statistical methods and implementing them practically.

Matrix Algebra in R

To refresh your matrix algebra skills with implementations in R, check out my [Crash Course on Matrix Algebra in R](#).

1.4.2 The list

In **lists**, you can organize different kinds of data. E.g., consider the following example:

```
mylist = list(
  "Some_Numbers" = c(66, 76, 55, 12, 4, 66, 8, 99),
  "Animals"       = c("Rabbit", "Cat", "Elephant"),
  "My_Series"     = c(30:1)
)
```

A very useful function to find specific values and entries within lists is the `str()`-function:

```
str(mylist)
```

```
List of 3
 $ Some_Numbers: num [1:8] 66 76 55 12 4 66 8 99
 $ Animals     : chr [1:3] "Rabbit" "Cat" "Elephant"
 $ My_Series   : int [1:30] 30 29 28 27 26 25 24 23 22 21 ...
```

1.4.3 The data frame

A `data.frame` is a `list` object with more formal restrictions (e.g., an equal number of rows for all columns). As indicated by its name, a `data.frame` object is designed to store data:

```
mydataframe = data.frame(  
  "Credit_Default" = c( 0, 0, 1, 0, 1, 1),  
  "Age"            = c(35,41,55,36,44,26),  
  "Loan_in_1000_EUR" = c(55,65,23,12,98,76)  
)
```

The `data()` command lists all sample data sets available in R. Let us have a look at the dataset `mtcars`. It is a `data.frame` object and contains data on several aspects of 32 automobiles from 1974.


`mtcars`

	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
Mazda RX4	21.0	6	160.0	110	3.90	2.620	16.46	0	1	4	4
Mazda RX4 Wag	21.0	6	160.0	110	3.90	2.875	17.02	0	1	4	4
Datsun 710	22.8	4	108.0	93	3.85	2.320	18.61	1	1	4	1
Hornet 4 Drive	21.4	6	258.0	110	3.08	3.215	19.44	1	0	3	1
Hornet Sportabout	18.7	8	360.0	175	3.15	3.440	17.02	0	0	3	2
Valiant	18.1	6	225.0	105	2.76	3.460	20.22	1	0	3	1
Duster 360	14.3	8	360.0	245	3.21	3.570	15.84	0	0	3	4
Merc 240D	24.4	4	146.7	62	3.69	3.190	20.00	1	0	4	2
Merc 230	22.8	4	140.8	95	3.92	3.150	22.90	1	0	4	2
Merc 280	19.2	6	167.6	123	3.92	3.440	18.30	1	0	4	4
Merc 280C	17.8	6	167.6	123	3.92	3.440	18.90	1	0	4	4
Merc 450SE	16.4	8	275.8	180	3.07	4.070	17.40	0	0	3	3
Merc 450SL	17.3	8	275.8	180	3.07	3.730	17.60	0	0	3	3
Merc 450SLC	15.2	8	275.8	180	3.07	3.780	18.00	0	0	3	3
Cadillac Fleetwood	10.4	8	472.0	205	2.93	5.250	17.98	0	0	3	4
Lincoln Continental	10.4	8	460.0	215	3.00	5.424	17.82	0	0	3	4
Chrysler Imperial	14.7	8	440.0	230	3.23	5.345	17.42	0	0	3	4
Fiat 128	32.4	4	78.7	66	4.08	2.200	19.47	1	1	4	1
Honda Civic	30.4	4	75.7	52	4.93	1.615	18.52	1	1	4	2
Toyota Corolla	33.9	4	71.1	65	4.22	1.835	19.90	1	1	4	1
Toyota Corona	21.5	4	120.1	97	3.70	2.465	20.01	1	0	3	1
Dodge Challenger	15.5	8	318.0	150	2.76	3.520	16.87	0	0	3	2
AMC Javelin	15.2	8	304.0	150	3.15	3.435	17.30	0	0	3	2
Camaro Z28	13.3	8	350.0	245	3.73	3.840	15.41	0	0	3	4

Pontiac Firebird	19.2	8	400.0	175	3.08	3.845	17.05	0	0	3	2
Fiat X1-9	27.3	4	79.0	66	4.08	1.935	18.90	1	1	4	1
Porsche 914-2	26.0	4	120.3	91	4.43	2.140	16.70	0	1	5	2
Lotus Europa	30.4	4	95.1	113	3.77	1.513	16.90	1	1	5	2
Ford Pantera L	15.8	8	351.0	264	4.22	3.170	14.50	0	1	5	4
Ferrari Dino	19.7	6	145.0	175	3.62	2.770	15.50	0	1	5	6
Maserati Bora	15.0	8	301.0	335	3.54	3.570	14.60	0	1	5	8
Volvo 142E	21.4	4	121.0	109	4.11	2.780	18.60	1	1	4	2

With the function `subset` we can select variables and subsets of a dataframe. Let's create a scatterplot of the variables `mpg` (miles per gallon) and `wt` weight (in 1000 lbs).

```
plot(subset(mtcars, select = c(wt, mpg)))
```



A `data.frame` is also useful in a time series context. Since time series data typically include a calendar date for each observation, the observation and date can be stored together as a `data.frame`. R provides the class `Date` for calendar dates, which can be generated with the function `as.Date()`.

```
d = as.Date("2021-04-01") # a data object to store dates
class(d) # to get the object class
```

```
[1] "Date"
```

```

myseries = c(16,17,18,16,15,19)
mydates = seq.Date(as.Date("2021-04-01"), by=1, length.out = 6)
mytimeseries = data.frame(mydates, myseries)
mytimeseries

```

	mydates	myseries
1	2021-04-01	16
2	2021-04-02	17
3	2021-04-03	18
4	2021-04-04	16
5	2021-04-05	15
6	2021-04-06	19

1.4.4 The ts object

A `ts` (time series) object is tailored explicitly to time series with a yearly time basis and an equidistant observation horizon, such as annual, quarterly, and monthly data. It assigns a specific year/quarter/month to each vector entry.

```

myts = ts(c(66, 76, 55, 12, 4, 66, 8, 99), start = 2020, frequency = 4)
myts

```

	Qtr1	Qtr2	Qtr3	Qtr4
2020	66	76	55	12
2021	4	66	8	99

```

anothertimeseries = ts(1:50, start = 2015, frequency = 12)
anothertimeseries

```

	Jan	Feb	Mar	Apr	May	Jun	Jul	Aug	Sep	Oct	Nov	Dec
2015	1	2	3	4	5	6	7	8	9	10	11	12
2016	13	14	15	16	17	18	19	20	21	22	23	24
2017	25	26	27	28	29	30	31	32	33	34	35	36
2018	37	38	39	40	41	42	43	44	45	46	47	48
2019	49	50										

```

# The window() command selects the time series observations for a given subperiod
window(anothertimeseries, start=2015.5, end=2017.5)

```

	Jan	Feb	Mar	Apr	May	Jun	Jul	Aug	Sep	Oct	Nov	Dec
2015							7	8	9	10	11	12
2016	13	14	15	16	17	18	19	20	21	22	23	24
2017	25	26	27	28	29	30	31					

The `data()` command lists all sample data sets available in R. Let us have a look at the dataset `AirPassengers`. It is a `ts` object and contains data on monthly totals of international airline passengers from 1949 to 1960.

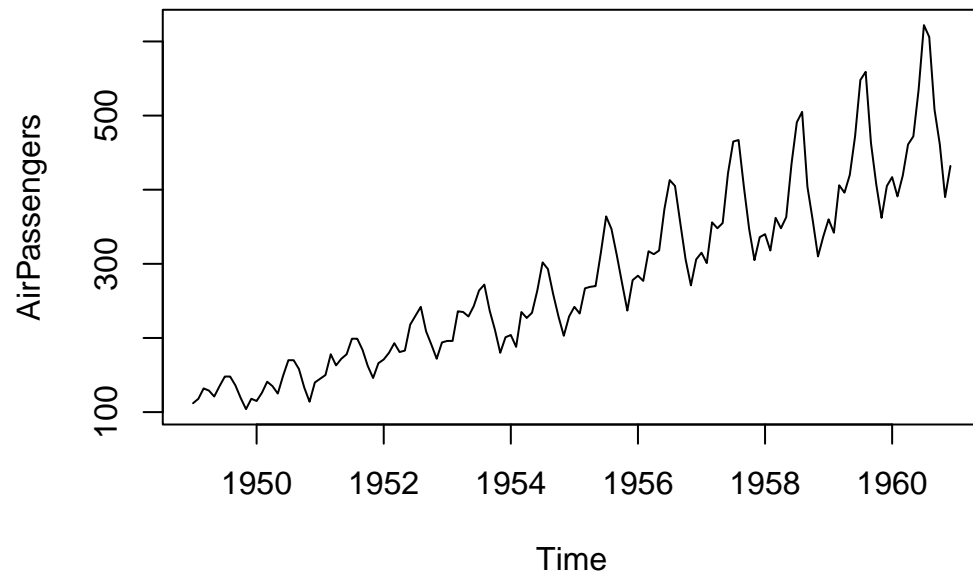
```
data() # lists all datasets currently loaded in the R environment
?AirPassengers # get more information about the dataset
AirPassengers
```

	Jan	Feb	Mar	Apr	May	Jun	Jul	Aug	Sep	Oct	Nov	Dec
1949	112	118	132	129	121	135	148	148	136	119	104	118
1950	115	126	141	135	125	149	170	170	158	133	114	140
1951	145	150	178	163	172	178	199	199	184	162	146	166
1952	171	180	193	181	183	218	230	242	209	191	172	194
1953	196	196	236	235	229	243	264	272	237	211	180	201
1954	204	188	235	227	234	264	302	293	259	229	203	229
1955	242	233	267	269	270	315	364	347	312	274	237	278
1956	284	277	317	313	318	374	413	405	355	306	271	306
1957	315	301	356	348	355	422	465	467	404	347	305	336
1958	340	318	362	348	363	435	491	505	404	359	310	337
1959	360	342	406	396	420	472	548	559	463	407	362	405
1960	417	391	419	461	472	535	622	606	508	461	390	432

```
class(AirPassengers) # AirPassengers is a ts object
```

```
[1] "ts"
```

```
plot(AirPassengers)
```



2 Packages

One of the best features of R are the large number of contributed packages from the statistical community. The list of all packages on CRAN is impressive! Take a look at it [here](#). You find R packages for almost any statistical method out there. Many statisticians provide R packages to accompany their research. Some packages also provide additional functionality for R or include datasets.

2.1 The `xts` package

Let us look at a time series specific package: the `xts` package. It can be installed using the `install.packages()` function.

```
install.packages("xts")
```

The `xts` package provides the class `xts`, which has certain advantages over `ts`. A `ts` object can specify the frequency of a time series only as a portion of a year (1 for yearly, 4 for quarterly, 12 for monthly data). This scheme is convenient for regular macroeconomic time series but impractical for daily data (leap year problem), high-frequency data, or irregularly collected data. In an `xts` object, we are much more flexible and manually assign a specific time index to each observation in the time series.

Once installed, the package only has to be loaded at the beginning of a new R session, which is done with the command `library(xts)`.


```
library(xts)
myts = ts(c(66, 76, 55, 12, 4, 66, 8, 99), start = 2020, frequency = 4)
as.xts(myts) # convert a ts object into an xts object
```

```
      [,1]
2020 Q1   66
2020 Q2   76
2020 Q3   55
2020 Q4   12
2021 Q1    4
```



```
2021 Q2    66
2021 Q3     8
2021 Q4    99
```

```
# we may assign irregular time points:
dates = seq.Date(as.Date("2023-01-01"), by = 7, length.out = 7)
dates[8] = as.Date("2023-03-01")
my.xts = xts(myts, dates)
plot(my.xts)
```



2.2 Data packages

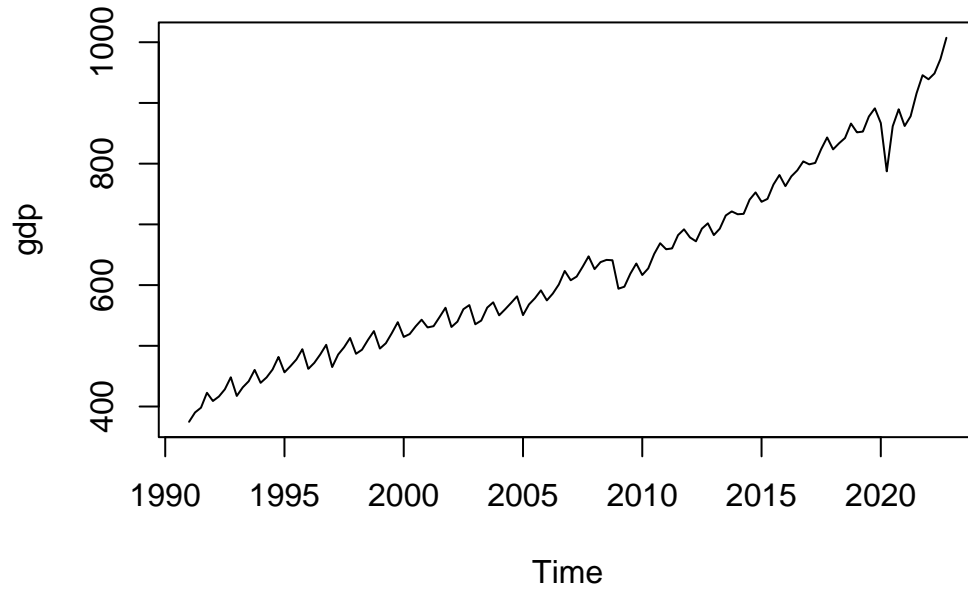
For teaching, I have created the package `teachingdata`, which contains some current datasets. The package is not available on CRAN (your package must meet specific quality standards and go through a review process to be accepted there), but I have created a GitHub repository to make it accessible. We need the package `remotes` and its function `install_github()` to install a package from a GitHub repository.

```
install.packages("remotes")
remotes::install_github("ottosven/teachingdata")
```

Let's have a closer look at the data from the `teachingdata` package.


```
library(teachingdata)
data(package = "teachingdata")
plot(gdp, main = "Quarterly GDP Germany")
```

Quarterly GDP Germany




```
plot(infl, main="Monthly CPI inflation rate Germany")
```

Monthly CPI inflation rate Germany

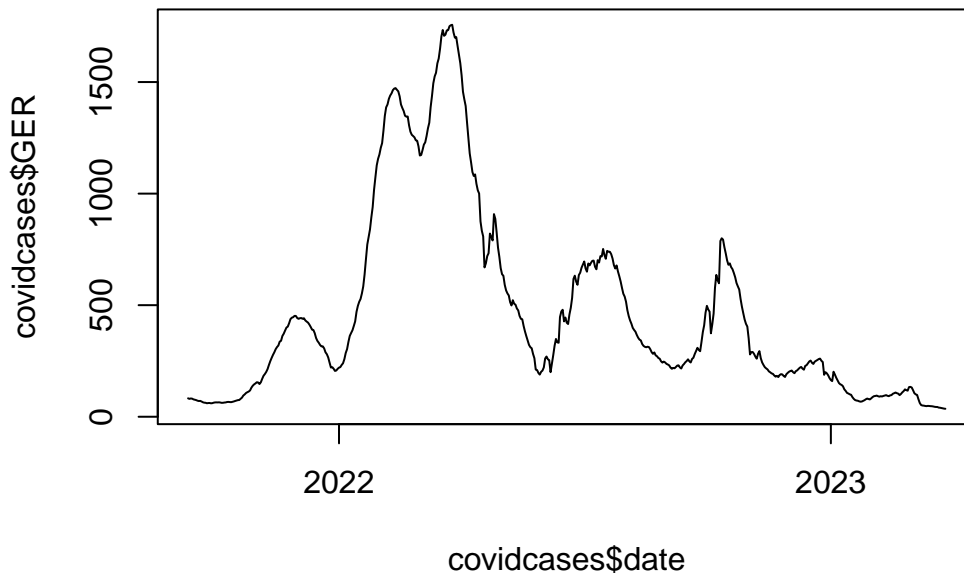


```
plot(temp, main="Average temperature Germany")
```



```
plot(covidcases$date, covidcases$GER, type="l",  
     main="Incidence number of reported Covid-19 infections Germany")
```

Incidence number of reported Covid-19 infections Germa



2.3 The tidyverse

The **tidyverse** is a collection of packages that lets you import, manipulate, explore, visualize, and model data in a harmonized and consistent way.

Installing the **tidyverse** package:

```
install.packages("tidyverse")
```

In this lecture, we will mainly use R to theoretically understand the learned statistical and econometric methods and apply them illustratively. For this purpose, base R is entirely sufficient. However, **tidyverse** has become state of the art for applied work with large data sets and is especially recommended for data management and visualization.

To give you a flavor of the tidyverse, let us briefly discuss the **ggplot2** and **tibble** packages, which are part of the **tidyverse**.


```
library(tidyverse)
```

Nice plots can be produced using the R-package **ggplot2**. Let's plot the **iris** dataset, which is contained in base R.

```
class(iris) # iris is a data.frame
```

```
[1] "data.frame"
```

```
iris |>  
  ggplot(aes(x = Sepal.Length, y = Petal.Length, color = Species)) +  
  geom_point()
```



A `data.frame` in the tidyverse is called `tibble`. A `tibble` is sometimes more flexible and convenient for manipulating and printing data. Let's transform the `iris` data frame into a `tibble`.

```
iris.tbl = as_tibble(iris)  
iris.tbl # iris.tbl is a tibble
```

```
# A tibble: 150 x 5
```

```
  Sepal.Length Sepal.Width Petal.Length Petal.Width Species  
    <dbl>         <dbl>         <dbl>         <dbl> <fct>
```

```

1      5.1      3.5      1.4      0.2 setosa
2      4.9      3      1.4      0.2 setosa
3      4.7      3.2      1.3      0.2 setosa
4      4.6      3.1      1.5      0.2 setosa
5      5        3.6      1.4      0.2 setosa
6      5.4      3.9      1.7      0.4 setosa
7      4.6      3.4      1.4      0.3 setosa
8      5        3.4      1.5      0.2 setosa
9      4.4      2.9      1.4      0.2 setosa
10     4.9      3.1      1.5      0.1 setosa
# i 140 more rows

```

As an extension, a `tsibble` object is a `tibble` with an additional time series structure. It contains a specific *index* variable corresponding to the observation's time index. Let us convert the `covidcases` data into a `tsibble`. To visualize a `tsibble` we also need the `fable` package.

```

library(tsibble)
library(fable)

```

In a `tsibble` object, we can define so-called *key* variables, which define the subjects or individuals measured over time. Key variables also allow easy processing of panel data in R.

In the `covidcases` example, the key variables are the federal states, and the time series is the incidence numbers. Since a simultaneous display of the curves of all federal states would produce a very cluttered plot, we select only the total Germany, Nordrhein-Westfalen, and Berlin. The different steps can be represented in tidyverse as a sequence of multiple operations using the pipe operator `|>` (other pipes like `%>%` do a similar job).

```

covid.tsibble = as_tsibble(covidcases, index=date) |>
  pivot_longer(-date, names_to = "state", values_to = "incidence") |>
  filter(state %in% c("GER", "NW", "BE"))
covid.tsibble

```

```

# A tsibble: 1,689 x 3 [1D]
# Key:      state [3]
   date      state incidence
  <date>    <chr>    <dbl>
1 2021-09-11 BE      83.5
2 2021-09-11 NW     103.
3 2021-09-11 GER     82.7
4 2021-09-12 BE     84.3

```

```


5 2021-09-12 NW      101.
6 2021-09-12 GER     80.1
7 2021-09-13 BE      83.7
8 2021-09-13 NW     99.3
9 2021-09-13 GER     81.8
10 2021-09-14 BE     84.9
# i 1,679 more rows

```

```

covid.tsibble |>
  autoplot(incidence) + theme(axis.title.x=element_blank())

```



For an introduction to the **tidyverse** and to learn more about the packages and functions used above, have a look at the book [R for Data Science](#). To learn more about visualizing and analyzing time series data using the **tsibble** and **fable** packages, I recommend the textbook [Forecasting: principles and practice](#).